

ML Assignment Report

Akash Bankar (23CS60R65)

[Google Drive Project Link](#)

1. Preprocessing Details

1. Used NLTK tweet tokenizer to tokenise emoji and urls
2. Used Emoji module to demojize emojis. i.e. Emoji to text
3. Removed Punctuations
4. Removed Stopwords
5. Used Porter stemmer to stemming
6. Used Tfidfvectorizer

2. Hyperparameter Tuning

Support Vector Machine (SVM)

- **Hyperparameters Tuned:** We tuned the following hyperparameters for SVM:
 - Kernel: ('linear', 'rbf', 'poly', 'sigmoid')
 - C: [1, 10, 100, 1000, 0.1]
 - Gamma: ('scale', 'auto')
- **Tuning Technique:** GridSearchCV was employed for hyperparameter tuning.
- **Best Attributes Obtained** - (C=10, kernel='sigmoid', gamma='auto')

Logistic Regression

- **Hyperparameters Tuned:** We explored the effect of different hyperparameters for Logistic Regression,
- **Tuning Technique:** Brute Force for each solver, penalty, C

- **Best Attributes Obtained** - (solver='saga', penalty='none')

KNN

- **Hyperparameters Tuned:** We tuned the following hyperparameters for SVM:

```
{'n_neighbors': [2, 3, 4, 5, 10, 15, 20, 25, 30],
 'weights' : ('uniform', 'distance'),
 'algorithm' : ('auto', 'ball_tree', 'kd_tree', 'brute')}
```

- **Tuning Technique:** GridSearchCV was employed for hyperparameter tuning.
- **Best Attributes Obtained** - (n_neighbors=4, weights='distance')

MLPClassifier (Neural Network)

- **Hyperparameters Tuned:** Tried to learn a lot of parameters as follows -

hidden_layer_sizes

activation{'identity', 'logistic', 'tanh', 'relu'}

solver{'lbfgs', 'sgd', 'adam'}

learning_rate{'constant', 'invscaling', 'adaptive'}

But each iteration was taking a lot of time to learn, so couldn't tune. Used defaults.

Code Example (whole ipynb attached) -

```
from sklearn.model_selection import GridSearchCV
parameters = {'kernel':('linear', 'rbf', 'poly', 'sigmoid'),
svc = SVC()
clf = GridSearchCV(svc, parameters)
clf.fit(X_train_text, y_train)
print(clf.cv_results_)
```

3. Results

KNN

- **Best Score:**

precision recall f1-score support

fake	0.88	0.92	0.90	501
real	0.93	0.88	0.90	559
accuracy			0.90	1060

macro avg 0.90 0.90 0.90 1060

weighted avg 0.90 0.90 0.90 1060

confusion matrix : [[462 39]

[66 493]]

accuracy score : 0.9009433962264151

recall_average : 0.8819320214669052

precision_average : 0.9266917293233082

f1_score : 0.9037580201649863

Support Vector Machine (SVM)

- **Best Score:**

precision recall f1-score support

fake	0.92	0.95	0.93	501
real	0.95	0.93	0.94	559
accuracy			0.94	1060

macro avg 0.94 0.94 0.94 1060

weighted avg 0.94 0.94 0.94 1060

confusion matrix : [[475 26]

[41 518]]

accuracy score : 0.9367924528301886

recall_average : 0.9266547406082289

precision_average : 0.9522058823529411

f1_score : 0.9392565729827741

Logistic Regression

- **Best Score:**

precision recall f1-score support

fake	0.92	0.95	0.93	501
real	0.96	0.92	0.94	559
accuracy			0.94	1060

macro avg 0.94 0.94 0.94 1060

weighted avg 0.94 0.94 0.94 1060

confusion matrix : [[477 24]

[44 515]]

accuracy score : 0.9358490566037736

recall_average : 0.9212880143112702

precision_average : 0.9554730983302412

f1_score : 0.9380692167577414

MLPClassifier (Neural Network)

- **Best Score:**

precision recall f1-score support

fake	0.92	0.96	0.94	501
real	0.96	0.92	0.94	559
accuracy			0.94	1060

macro avg 0.94 0.94 0.94 1060

weighted avg 0.94 0.94 0.94 1060

confusion matrix : [[480 21]

[44 515]]

accuracy score : 0.9386792452830188
recall_average : 0.9212880143112702
precision_average : 0.960820895522388
f1_score : 0.9406392694063926

KMeans

- **Best Score:**

precision recall f1-score support

fake	0.04	0.01	0.02	501
real	0.45	0.73	0.56	559
accuracy			0.39	1060

macro avg 0.24 0.37 0.29 1060
weighted avg 0.26 0.39 0.30 1060

confusion matrix : [[6 495]
[152 407]]

accuracy score : 0.389622641509434
recall_average : 0.7280858676207513
precision_average : 0.45121951219512196
f1_score : 0.5571526351813826

FastText

- **Best Score:**

Precision - 0.9837264150943397

Recall - 0.9837264150943397